




Distilling Human Knowledge


<https://likert.net>

Pitch Deck

How to train a model



Training process correlates with model intelligence



RLHF is an emerging market

- optimization targets are **data volume & quality**
- training methodology is mostly constant
- bottleneck metrics are **average handle time (AHT), researcher-labeler alignment.**
- data is expensive!

15x

YoY Headcount of
labelers @ OpenAI

80\$

Average hourly wage

30,000

Currently employed

What about data vendors?

Labs don't like working with vendors:

- very costly
- slower iterations
- lower quality
- researcher-labeler alignment drifts
- trust issues (eg Scale's acq. by Meta)

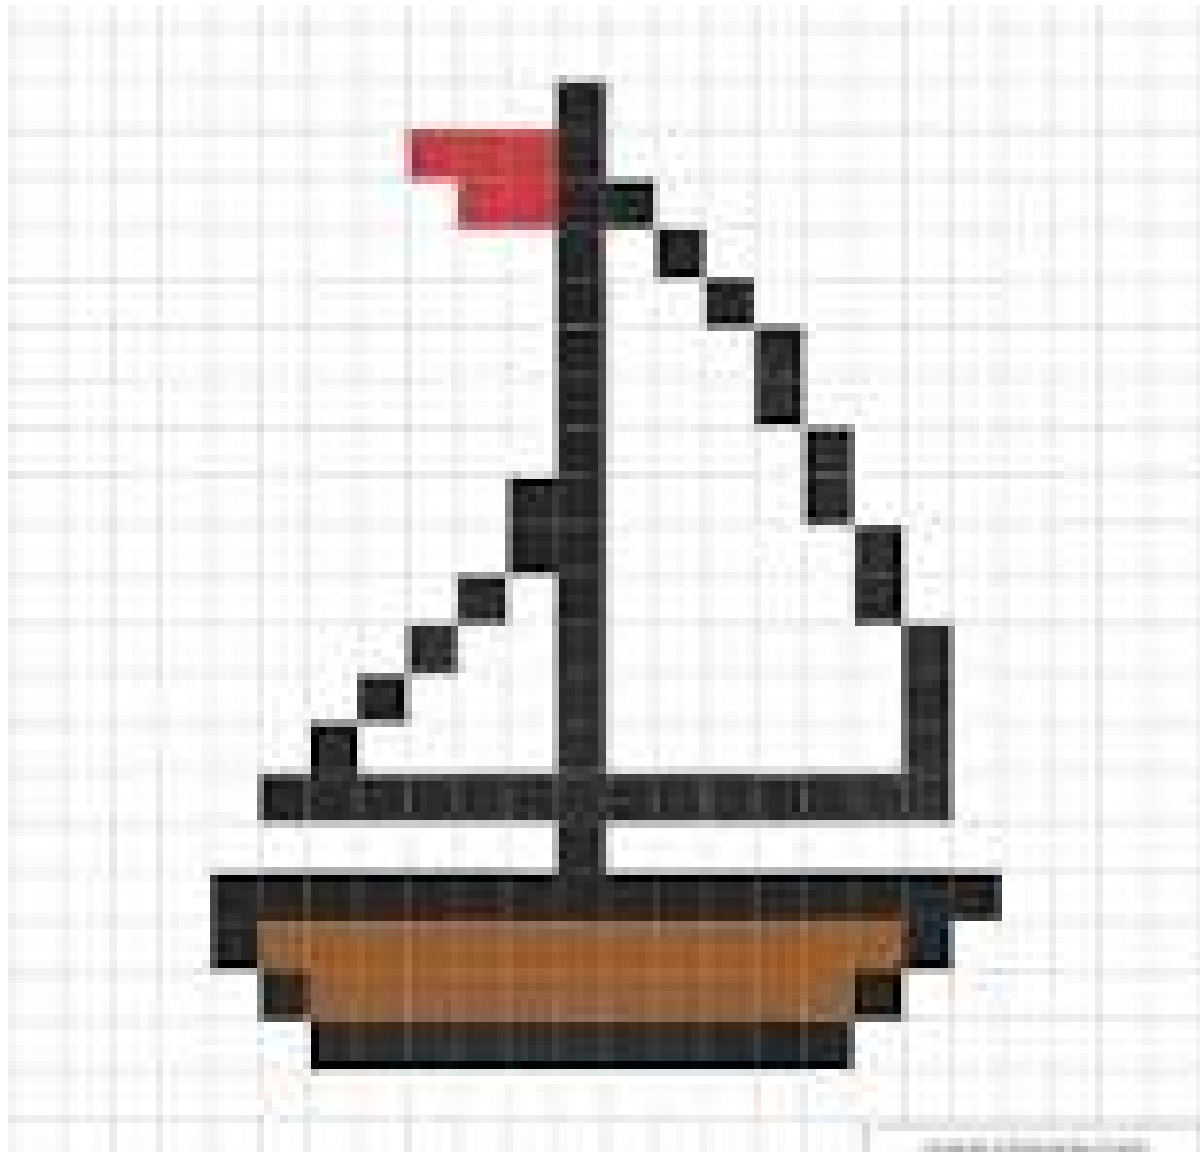
They're just the only solution available right now; Likert wants to change that.

Example RLHF Task

Rank your preferences 1st, 2nd, 3rd;
Afterwards, write a paragraph of
feedback for each completion.

Prompt:
“Generate an oil painting of a sailboat on a
river besides some houses in the style of
Claude Monet”.

Completion A



(bad: wrong medium)

Completion B



(bad: steamboat)

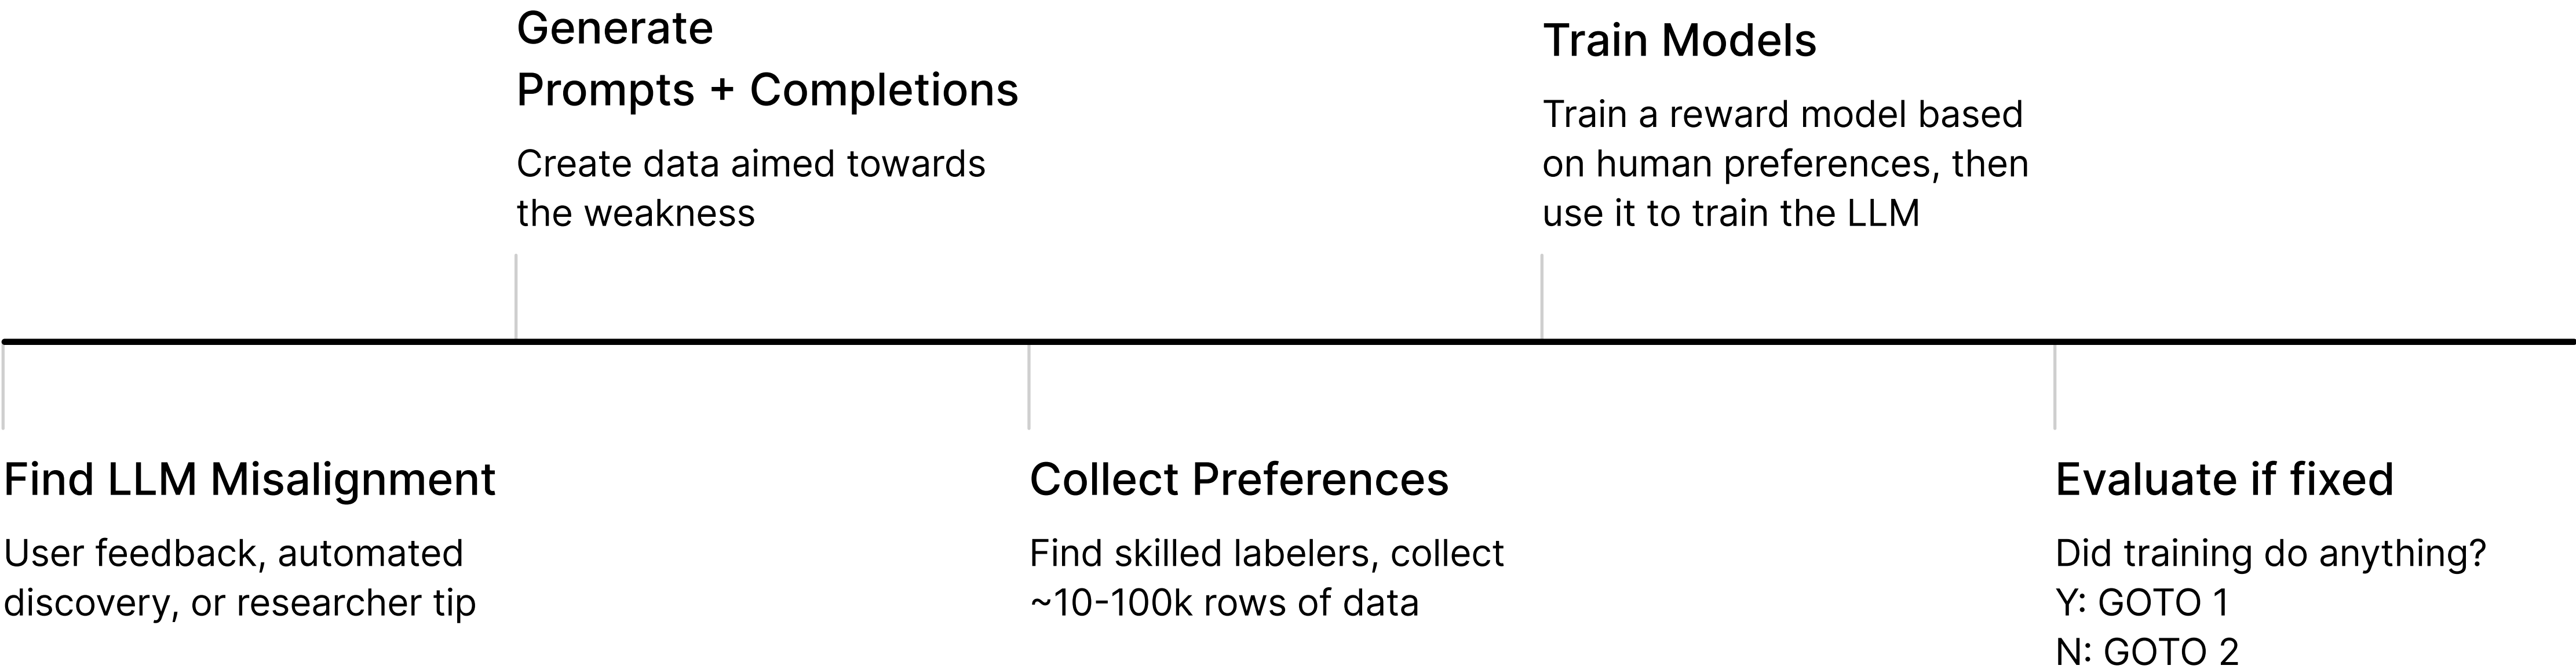
Completion C



(what we want)

Concretely, Likert

What Likert Automates



Market

Under \$1B
42 LABS



\$1B-\$3B
30 LABS



\$3B-\$10B
23 LABS



Over \$10B
10 LABS

